

# Zero-Shot Radiological Screening with Contrastive Pretraining

*MedVista Diagnostics -- Building a Language-Guided Medical Image Retrieval System*

---

---

## Section 1: Industry Context and Business Problem

### Industry: Healthcare -- Diagnostic Radiology

MedVista Diagnostics is a mid-size healthcare AI company serving 340 hospitals across the southeastern United States. Their core product is an AI-assisted radiology workflow system that helps radiologists triage and interpret medical images more efficiently.

### Company Profile

- **Founded:** 2019
- **Employees:** 180 (45 ML engineers, 12 radiologists on advisory board)
- **Revenue:** USD 28M ARR from SaaS contracts with hospital networks
- **Flagship Product:** RadAssist -- an AI copilot for radiologists that flags potential abnormalities in chest X-rays, CT scans, and MRIs

### The Business Challenge

MedVista's current system is built on traditional supervised classification models. Each model is trained on a specific pathology: one model for pneumonia detection, another for cardiomegaly, another for pleural effusion, and so on. This creates three critical problems:

1. **Label Bottleneck:** Training each new pathology detector requires 10,000--50,000 labeled images, annotated by board-certified radiologists at USD 15--40 per annotation. Adding support for a new condition costs USD 150K--500K in labeling alone.
2. **Long Tail of Conditions:** There are over 200 clinically significant radiological findings. MedVista currently supports only 14. Rare conditions (e.g., sarcoidosis, pneumomediastinum) will never have enough labeled data for supervised training.
3. **Deployment Rigidity:** When a new pathology needs to be detected, MedVista must retrain a new model, validate it against FDA requirements, and deploy it -- a process that takes 6--9 months. Competitors with more flexible systems are winning contracts.

## Stakes

Three major hospital networks (representing USD 8.2M in annual revenue) have issued RFPs requiring support for at least 50 radiological findings. MedVista's current 14-condition system will lose these contracts unless they dramatically expand coverage.

## Constraints

- Must operate on standard radiology PACS (Picture Archiving and Communication System) hardware
- Inference latency must be under 500ms per image
- Must maintain or exceed current per-pathology AUC (>0.88) for the 14 existing conditions
- FDA 510(k) clearance pathway requires interpretable similarity scores, not black-box classifications
- Training data is limited: USD 200K budget for additional annotations

## Proposed Solution

Build a contrastive pretraining system (CLIP-style) that learns a shared embedding space between radiology images and clinical text descriptions. This enables:

- **Zero-shot classification:** Detect any condition by providing its text description, no labeled data needed
  - **Text-guided retrieval:** Find similar cases by describing symptoms in natural language
  - **Flexible expansion:** Add new conditions in minutes (write a text prompt) instead of months (train a new model)
- 

## Section 2: Technical Problem Formulation

### Problem Type: Multimodal Contrastive Representation Learning

**Justification:** The core challenge is learning an alignment between visual features in medical images and semantic concepts expressed in clinical text. This is fundamentally a representation learning problem where the goal is to create a shared embedding space that captures medically meaningful relationships between images and their textual descriptions. Contrastive pretraining is the appropriate approach because:

1. It does not require per-pathology labeled data
2. It leverages existing radiology reports as natural language supervision
3. It enables zero-shot generalization to unseen conditions
4. The learned embeddings are inherently interpretable (similar images have similar embeddings)

## Input/Output Specifications

**Inputs:** - Medical images: chest X-rays, 224x224 grayscale, normalized to [0, 1] - Clinical text: radiology report findings, tokenized to max 77 tokens

**Outputs:** - Image embeddings:  $z_I \in \mathbb{R}^{512}$ , L2-normalized - Text embeddings:  $z_T \in \mathbb{R}^{512}$ , L2-normalized - Similarity scores:  $s_{ij} = z_{I_i} \cdot z_{T_j} \in [-1, 1]$

## Mathematical Foundation

The model consists of two encoders:

$$f_\theta : \mathcal{X}_{\text{image}} \rightarrow \mathbb{R}^d, \quad g_\phi : \mathcal{X}_{\text{text}} \rightarrow \mathbb{R}^d$$

where  $d = 512$  is the embedding dimension.

For a batch of  $N$  image-text pairs  $\{(x_{I_i}, x_{T_i})\}_{i=1}^N$ , we compute normalized embeddings:

$$\hat{z}_{I_i} = \frac{f_\theta(x_{I_i})}{\|f_\theta(x_{I_i})\|_2}, \quad \hat{z}_{T_i} = \frac{g_\phi(x_{T_i})}{\|g_\phi(x_{T_i})\|_2}$$

## Loss Function

The loss is a symmetric InfoNCE objective:

$$\mathcal{L} = \frac{1}{2} (\mathcal{L}_{I \rightarrow T} + \mathcal{L}_{T \rightarrow I})$$

where:

$$\mathcal{L}_{I \rightarrow T} = -\frac{1}{N} \sum_{i=1}^N \log \frac{\exp(\hat{z}_{I_i} \cdot \hat{z}_{T_i} / \tau)}{\sum_{j=1}^N \exp(\hat{z}_{I_i} \cdot \hat{z}_{T_j} / \tau)}$$

### Per-term justification:

- **Numerator**  $\exp(\hat{z}_{I_i} \cdot \hat{z}_{T_i} / \tau)$ : Measures alignment between matched image-text pairs. Maximizing this pushes correct pairs together in embedding space.
- **Denominator**  $\sum_j \exp(\hat{z}_{I_i} \cdot \hat{z}_{T_j} / \tau)$ : Normalizes across all texts in the batch, creating an implicit  $N$ -way classification problem. This pushes non-matching pairs apart.
- **Temperature**  $\tau$ : Controls the concentration of the distribution. Learned during training. Lower values increase discrimination but risk training instability.
- **Symmetry**  $\frac{1}{2}(\mathcal{L}_{I \rightarrow T} + \mathcal{L}_{T \rightarrow I})$ : Ensures both modalities contribute equally. Without symmetry, the model might learn good image representations but poor text representations (or vice versa).
- **Log and negative sign**: Converting the softmax probability to negative log-likelihood creates a proper cross-entropy loss that can be minimized with gradient descent.

## Evaluation Metrics

Metric	Description	Target
Zero-shot AUC	Per-pathology AUC using text prompts only	>0.80
Retrieval R@5	Fraction of queries where correct match is in top 5	>0.70
Embedding alignment	Mean cosine similarity for matched pairs	>0.60
Inference latency	Time per image embedding	<200ms
Coverage	Number of conditions detectable without fine-tuning	>50

## Baseline

The baseline is MedVista's existing supervised ResNet-50 classifiers, each trained on 10K+ labeled images per pathology. These achieve AUC >0.88 on their 14 supported conditions but cannot generalize to unseen conditions.

## Why Contrastive Pretraining

Supervised classification requires expensive per-class labels and cannot scale to hundreds of conditions. Contrastive pretraining solves this by:

1. Learning from radiology reports (free, existing data) rather than manual labels
2. Enabling zero-shot transfer to any condition describable in text
3. Producing interpretable similarity scores (required for FDA clearance)
4. Supporting both classification and retrieval in a unified system

---

## Section 3: Implementation Notebook Structure

---

### 3.1 Data Loading and Preparation

```
def load_medical_dataset(data_dir, split='train'):
    """
    Load the chest X-ray dataset with paired radiology reports.

    Args:
        data_dir: path to dataset root
        split: 'train', 'val', or 'test'

    Returns:
        images: list of image paths
        reports: list of radiology report strings
        labels: dict mapping pathology names to binary labels

    # ===== TODO =====
    # Step 1: Load image paths from data_dir/split/
    # Step 2: Load corresponding radiology reports from CSV
    # Step 3: Parse multi-label annotations for evaluation
    # Step 4: Apply preprocessing (resize to 224x224, normalize)
    # Hint: Use torchvision.transforms for image preprocessing
    # =====
    """
```

```
"""
pass
```

## 3.2 Exploratory Data Analysis

```
def exploratory_analysis(images, reports, labels):
    """
    Analyze the dataset distribution and characteristics.

    Tasks:
    - Plot label frequency distribution across pathologies
    - Show sample images with their report excerpts
    - Compute report length statistics (tokens, sentences)
    - Identify class imbalance and rare conditions
    - Visualize image quality distribution (brightness, contrast)

    # ===== TODO =====
    # Step 1: Count frequency of each pathology label
    # Step 2: Plot horizontal bar chart of label frequencies
    # Step 3: Display 3x3 grid of sample images with report snippets
    # Step 4: Histogram of report lengths
    # Step 5: Flag conditions with fewer than 100 examples
    # =====
    """
    pass
```

## 3.3 Baseline Model

```
def train_supervised_baseline(train_loader, val_loader, num_classes=14):
    """
    Train a supervised ResNet-50 classifier as baseline.

    Architecture:
    - ResNet-50 pretrained on ImageNet
    - Replace final FC layer with num_classes outputs
    - Binary cross-entropy loss for multi-label classification

    # ===== TODO =====
    # Step 1: Load pretrained ResNet-50
    # Step 2: Replace classifier head
    # Step 3: Train for 10 epochs with Adam optimizer
    # Step 4: Evaluate per-class AUC on validation set
    # Step 5: Report which conditions have AUC < 0.80
    # =====
    """
    pass
```

## 3.4 CLIP-Style Model Architecture

```
class MedCLIP(nn.Module):
    """
    Medical CLIP model with image and text encoders.

    Image encoder: ViT-B/16 initialized from ImageNet pretraining
    Text encoder: BioClinicalBERT (domain-specific)
    Projection: Linear layers mapping to 512-d shared space

    # ===== TODO =====
    # Step 1: Initialize image encoder (ViT-B/16 or ResNet-50)
    # Step 2: Initialize text encoder (BioClinicalBERT)
    # Step 3: Add projection heads (image_dim -> 512, text_dim -> 512)
    # Step 4: Initialize learnable temperature parameter
    # Step 5: Implement forward() that returns normalized embeddings
    # =====
    """

    def __init__(self, image_encoder, text_encoder, embed_dim=512):
        super().__init__()
        pass
```

```

def forward(self, images, input_ids, attention_mask):
    pass

def compute_loss(self, image_emb, text_emb):
    pass

```

## 3.5 Training Pipeline

```

def train_medclip(model, train_loader, val_loader, num_epochs=50):
    """
    Train the MedCLIP model with contrastive learning.

    Key considerations:
    - Use mixed-precision training (AMP) for memory efficiency
    - Gradient accumulation for effective batch size of 1024
    - Learning rate warmup + cosine decay schedule
    - Monitor contrastive accuracy and retrieval metrics

    # ===== TODO =====
    # Step 1: Set up optimizer (AdamW, lr=5e-5, weight_decay=0.01)
    # Step 2: Set up scheduler (linear warmup + cosine decay)
    # Step 3: Training loop with gradient accumulation
    # Step 4: Log contrastive loss, accuracy, and temperature
    # Step 5: Validate every epoch with retrieval metrics
    # =====
    """
    pass

```

## 3.6 Evaluation and Zero-Shot Classification

```

def evaluate_zero_shot(model, test_loader, pathology_names, templates):
    """
    Evaluate zero-shot classification performance.

    For each pathology:
    1. Create text prompts using templates
    2. Compute average text embedding (prompt ensembling)
    3. Compute image-text similarities for all test images
    4. Compute AUC using similarities as prediction scores

    # ===== TODO =====
    # Step 1: Pre-compute text embeddings for all pathologies
    # Step 2: For each test image, compute similarity to all pathologies
    # Step 3: Compute per-pathology AUC
    # Step 4: Compare with supervised baseline AUCs
    # Step 5: Identify conditions where zero-shot matches/exceeds supervised
    # =====
    """
    pass

```

## 3.7 Error Analysis

```

def error_analysis(model, test_loader, pathology_names):
    """
    Analyze failure modes of the zero-shot system.

    Tasks:
    - Compute confusion between similar conditions
    - Identify images where zero-shot fails but supervised succeeds
    - Analyze correlation between report quality and embedding quality
    - Test sensitivity to prompt wording variations
    - Examine embedding space for cluster quality

    # ===== TODO =====
    # Step 1: Collect all predictions and ground truth
    # Step 2: Build confusion matrix for top-3 predictions
    # Step 3: Visualize failure cases with images and similarity scores
    # Step 4: Test 5 different prompt templates per condition
    # Step 5: t-SNE of image embeddings colored by pathology

```

```
# =====
"""
pass
```

## 3.8 Deployment Considerations

```
def optimize_for_deployment(model, sample_input):
    """
    Prepare the model for production deployment.

    Tasks:
    - Export to ONNX format for cross-platform inference
    - Benchmark inference latency on CPU and GPU
    - Implement batch processing for throughput optimization
    - Set up embedding caching for text prompts
    - Validate numerical consistency between PyTorch and ONNX

    # ===== TODO =====
    # Step 1: Trace model with torch.jit.trace
    # Step 2: Export to ONNX with dynamic batch size
    # Step 3: Benchmark latency for single image and batch of 16
    # Step 4: Implement text embedding cache (dictionary)
    # Step 5: Verify ONNX output matches PyTorch within tolerance
    # =====
    """
    pass
```

## 3.9 Ethical Considerations

```
def fairness_analysis(model, test_loader, demographic_metadata):
    """
    Evaluate model fairness across demographic groups.

    Tasks:
    - Compute zero-shot AUC stratified by age, sex, and ethnicity
    - Identify conditions with significant performance gaps
    - Test for bias in text encoder (prompt sensitivity analysis)
    - Document limitations for FDA submission

    # ===== TODO =====
    # Step 1: Stratify test set by demographic attributes
    # Step 2: Compute per-group AUC for each pathology
    # Step 3: Flag conditions with AUC gap > 0.05 between groups
    # Step 4: Test if different prompts affect different groups differently
    # Step 5: Generate fairness report table
    # =====
    """
    pass
```

# Section 4: Production and System Design Extension

## System Architecture

The production MedCLIP system consists of the following components:

- 1. Image Ingestion Service:** Receives DICOM images from hospital PACS via HL7 FHIR interface, performs DICOM-to-PNG conversion, and applies standardized preprocessing (resizing, windowing, normalization).

- 2. Embedding Service:** GPU-backed microservice running the MedCLIP image encoder. Receives preprocessed images and returns 512-d normalized embeddings. Deployed as a containerized service with auto-scaling.
- 3. Prompt Registry:** Database of validated text prompts for each pathology, stored as pre-computed text embeddings. Updated when new conditions are added (no model retraining required).
- 4. Similarity Engine:** Computes cosine similarity between image embeddings and all registered pathology embeddings. Returns ranked results with confidence scores.
- 5. Result Aggregation:** Combines similarity scores with clinical rules (age-appropriate screening, anatomical constraints) to produce final findings.

## API Design

```

POST /api/v1/screen
Input: DICOM image (binary)
Output: {
  "findings": [
    {"pathology": "cardiomegaly", "confidence": 0.92, "rank": 1},
    {"pathology": "pleural_effusion", "confidence": 0.78, "rank": 2}
  ],
  "embedding": [0.12, -0.34, ...], // 512-d vector
  "latency_ms": 187,
  "model_version": "medclip-v2.1"
}

POST /api/v1/retrieve
Input: {"query": "bilateral ground-glass opacities", "top_k": 10}
Output: {
  "results": [
    {"image_id": "CXR-2024-00142", "similarity": 0.89, "metadata": {...}},
    ...
  ]
}

POST /api/v1/register-condition
Input: {"name": "sarcoidosis", "prompts": ["bilateral hilar lymphadenopathy", ...]}
Output: {"status": "registered", "embedding_id": "emb-00234"}

```

## Serving Infrastructure

- **GPU Cluster:** 4x NVIDIA A10G GPUs behind a load balancer
- **Batch Inference:** Images accumulated into batches of 16 for throughput optimization
- **Embedding Cache:** Redis-backed cache for text embeddings (invalidated on prompt updates)
- **Model Serving:** TorchServe with ONNX Runtime backend for inference optimization
- **Target Throughput:** 200 images/second at p99 latency <500ms

## Monitoring

Metric	Description	Alert Threshold
Inference latency (p95)	Image encoding time	>400ms

Metric	Description	Alert Threshold
Similarity distribution	Mean similarity for top prediction	Shift >0.05 from baseline
Embedding drift	Cosine distance between rolling averages	>0.10 per week
Zero-shot AUC	Weekly AUC on validation holdout	Drop >0.03
Cache hit rate	Fraction of text embeddings served from cache	<0.95

## Drift Detection

The system monitors for two types of drift:

- 1. Data drift:** The distribution of incoming images changes (new scanner types, different patient demographics, acquisition protocol changes). Detected by monitoring the mean and variance of image embeddings over time.
- 2. Concept drift:** The relationship between images and pathologies changes (new clinical guidelines, updated terminology). Detected by periodically evaluating zero-shot AUC on a held-out validation set.

When drift is detected, the system:

- Triggers an alert to the ML ops team
- Automatically increases the validation frequency
- Logs affected image embeddings for review
- Does NOT automatically retrain (regulatory constraint -- any model update requires FDA re-evaluation)

## A/B Testing

For evaluating model improvements:

- Shadow mode: new model runs in parallel, predictions logged but not surfaced
- Clinical validation: 2-week parallel run with radiologist review of discordant predictions
- Rollout: gradual traffic shift (10% -> 25% -> 50% -> 100%) over 4 weeks
- Rollback criterion: if per-pathology AUC drops by >0.02 for any existing condition

## CI/CD Pipeline

```

Code Push -> Unit Tests -> Integration Tests (synthetic data)
-> Model Training (full pipeline, ~8 hours)
-> Validation Suite (zero-shot AUC, retrieval R@5, fairness checks)
-> Staging Deployment (shadow mode, 48 hours)
-> Clinical Review (radiologist spot-check)
-> Production Deployment (canary -> full rollout)

```

## Cost Analysis

Component	Monthly Cost
GPU inference (4x A10G)	USD3,200
Storage (embeddings, images)	USD450
Redis cache	USD180
Monitoring and logging	USD120

<b>Component</b>	<b>Monthly Cost</b>
Training (periodic, 8x A100 for 12 hours)	USD960/run
<b>Total operational</b>	<b>USD4,910/month</b>

Compared to the supervised approach (training and maintaining 50+ separate models at ~USD2,000/model/year for annotation and retraining), MedCLIP reduces per-condition costs by approximately 85%.