

Case Study: Adaptive Content Recommendation with Policy Gradient Methods

NovaMind AI -- Personalizing Learning Paths in Real-Time

Section 1: Industry Context and Business Problem

Industry: EdTech -- Adaptive Learning Platforms

The global edtech market reached \\$340 billion in 2025, with adaptive learning platforms representing the fastest-growing segment. Unlike static content delivery, adaptive systems observe learner behavior in real time and adjust the sequence and difficulty of presented material to maximize learning outcomes.

Company Profile: NovaMind AI

NovaMind AI is a Series B edtech startup (\\$45M raised) headquartered in San Francisco. They operate an adaptive learning platform serving 2.3 million users across K-12, university, and corporate training verticals. Their platform delivers microlearning modules -- short, focused lessons of 3-15 minutes each.

Revenue model: \\$18/month per individual user, \\$12/month per enterprise seat (minimum 500 seats).

Current annual recurring revenue (ARR): \\$31M with 38% year-over-year growth.

The Business Challenge

NovaMind's current content recommendation system uses a rule-based approach: a fixed decision tree that assigns content based on pre-test scores and completion rates. This system has three critical problems:

- 1. One-size-fits-all sequencing.** A learner who scores 70% on a pre-test gets the same content sequence as every other learner with a 70% score, regardless of their learning pace, preferred content type, or knowledge gaps.
- 2. Delayed adaptation.** The system only re-evaluates content assignments after a learner completes a full module (10-15 minutes). Learners who are struggling receive no intervention during the module itself.

3. No exploration. The system never tries novel content sequences. It always recommends the "safe" path, missing opportunities to discover that some learners benefit from advanced content earlier or from revisiting fundamentals in a different order.

Business impact: - Learner completion rates have stagnated at 34% - Average session duration is declining (from 23 minutes to 17 minutes over 6 months) - Enterprise churn rate is 8.2% quarterly (industry benchmark: 5%) - NPS score dropped from 42 to 31

Stakes and Constraints

NovaMind must improve completion rates to 50% within 6 months to retain their largest enterprise client (\\$4.2M annual contract) and hit their Series C growth targets.

Technical constraints: - Recommendations must be generated in under 200ms - The system must work for new users with no history (cold start) - Content library: 12,000 microlearning modules across 840 topics - User interactions: ~8M per day - Model must be explainable for enterprise compliance

Why Policy Gradient Methods?

This problem is fundamentally a **sequential decision-making** problem under uncertainty. At each step, the system must choose which content to show next based on the learner's current state. The "reward" is delayed -- we only know if the content sequence was effective after the learner completes (or abandons) their session.

Value-based methods (like DQN) struggle here because: - The action space is large (12,000 modules to choose from) - Actions need a probability distribution for exploration, not a single best action - The policy must naturally handle stochastic exploration

Policy gradient methods are the right solution because they directly output a probability distribution over the 12,000 modules, naturally handle the large action space, and allow controlled exploration through the softmax temperature.

Section 2: Technical Problem Formulation

Problem Type: Sequential Recommendation as a Markov Decision Process

We frame the adaptive content recommendation problem as an MDP where: - **States** represent the learner's current knowledge state - **Actions** represent content module recommendations - **Rewards** capture learning effectiveness - **Transitions** model how the learner's state changes after interacting with content

Input/Output Specification

State vector $s_t \in \mathbb{R}^{256}$: - Learner embedding (128-dim): encodes historical performance, content preferences, and learning velocity - Session context (64-dim): current session duration, number of modules completed, time of day, device type - Knowledge state (64-dim): estimated mastery across 64 topic clusters

Action $a_t \in \{1, 2, \dots, 12000\}$: - Index of the microlearning module to recommend next

Reward r_t : - Module completion: +1.0 if the learner completes the module, 0.0 otherwise - Quiz performance: $+0.5 * (\text{quiz score} / 100)$ if the module contains a quiz - Engagement bonus: +0.3 if the learner spends at least 80% of expected time on the module - Negative signal: -0.5 if the learner skips the module within 30 seconds

Mathematical Foundation

Policy parameterization:

We use a neural network policy $\pi_\theta(a|s)$ that outputs a probability distribution over all 12,000 modules. However, directly computing softmax over 12,000 actions is computationally expensive. We use a two-stage approach:

Stage 1 (Candidate generation): A lightweight model retrieves the top-K=100 candidate modules based on topic relevance and difficulty matching.

Stage 2 (Policy scoring): The policy network π_θ scores only the K=100 candidates:

$$\pi_\theta(a_i|s) = \frac{e^{h_\theta(s, a_i)}}{\sum_{j=1}^K e^{h_\theta(s, a_j)}}$$

where $h_\theta(s, a_i)$ is the preference score for module a_i given state s .

Let us plug in simple numbers. Suppose K=3 candidate modules with scores $h(a_1) = 3.2$, $h(a_2) = 2.8$, $h(a_3) = 1.5$:

$$e^{3.2} = 24.53, \quad e^{2.8} = 16.44, \quad e^{1.5} = 4.48$$

sum = 45.45

$$\pi(a_1) = 0.54, \quad \pi(a_2) = 0.36, \quad \pi(a_3) = 0.10$$

Module a_1 has the highest probability but a_2 and a_3 still have non-trivial probability, allowing exploration.

Performance measure:

$$J(\theta) = \mathbb{E}_{\tau \sim \pi_\theta} \left[\sum_{t=0}^{T-1} \gamma^t r_t \right]$$

where T is the session length (variable, up to 20 interactions) and $\gamma = 0.95$.

Policy gradient update with baseline:

$$\nabla_{\theta} J(\theta) = \mathbb{E}_{\tau} \left[\sum_{t=0}^{T-1} \nabla_{\theta} \log \pi_{\theta}(a_t | s_t) \cdot (G_t - V_{\phi}(s_t)) \right]$$

We use an Actor-Critic architecture where: - **Actor** (π_{θ}): Recommends content - **Critic** (V_{ϕ}): Estimates expected session return

Loss Function

The total loss has three terms:

$$\mathcal{L} = \mathcal{L}_{\text{actor}} + c_1 \mathcal{L}_{\text{critic}} + c_2 \mathcal{L}_{\text{entropy}}$$

Term 1 -- Actor loss: Negative of the policy gradient objective

$$\mathcal{L}_{\text{actor}} = -\frac{1}{N} \sum_{i=1}^N \sum_{t=0}^{T_i-1} \log \pi_{\theta}(a_t^{(i)} | s_t^{(i)}) \cdot A_t^{(i)}$$

Justification: This directly optimizes the policy to increase the probability of actions that led to higher-than-expected returns and decrease the probability of those that led to lower-than-expected returns.

Term 2 -- Critic loss: Mean squared error between predicted and actual returns

$$\mathcal{L}_{\text{critic}} = \frac{1}{N} \sum_{i=1}^N \sum_{t=0}^{T_i-1} (V_{\phi}(s_t^{(i)}) - G_t^{(i)})^2$$

Justification: Training the critic to accurately predict returns provides a low-variance baseline for the actor, reducing gradient noise.

Term 3 -- Entropy bonus:

$$\mathcal{L}_{\text{entropy}} = -\frac{1}{N} \sum_{i=1}^N \sum_{t=0}^{T_i-1} H(\pi_{\theta}(\cdot | s_t^{(i)}))$$

Justification: Prevents the policy from collapsing to always recommending the same few modules. Controlled exploration is critical in an educational setting where diverse content exposure benefits learning.

Coefficients: $c_1 = 0.5$, $c_2 = 0.01$.

Evaluation Metrics

Metric	Definition	Target
Completion Rate	% of recommended modules fully completed	> 50%
Session Return	Average cumulative reward per session	> 8.0

Metric	Definition	Target
Knowledge Gain	Pre-post test score improvement	> 15 pp
Exploration Ratio	% of unique modules recommended in 1000 sessions	> 40%

Baseline Method

Rule-based recommender: The current production system that assigns content based on pre-test scores using a fixed decision tree. This achieves 34% completion rate and 5.2 average session return.

Section 3: Implementation Notebook Structure

3.1 Data Generation and Preprocessing

```
# TODO: Implement the data generator for simulated learner interactions

class LearnerEnvironment:
    """
        Simulated environment for adaptive content recommendation.
        Models learner behavior as an MDP.
    """
    def __init__(self, num_modules=100, state_dim=32, num_topics=10):
        """
            Args:
                num_modules: number of content modules available
                state_dim: dimension of the state vector
                num_topics: number of topic clusters
        """
        # ===== TODO =====
        # Initialize the environment:
        # 1. Create module difficulty levels (uniform 0-1)
        # 2. Create module-topic assignments
        # 3. Initialize learner knowledge state
        # =====
        pass

    def reset(self):
        """
            Reset the environment for a new session.
        """
        # ===== TODO =====
        # Return initial state vector
        # =====
        pass

    def step(self, action):
        """
            Execute one step: learner interacts with recommended module.

            Returns: next_state, reward, done, info
        """
        # ===== TODO =====
        # 1. Check if module difficulty matches learner level
        # 2. Compute completion probability
        # 3. Compute reward based on completion and engagement
        # 4. Update learner knowledge state
        # 5. Return (next_state, reward, done, info)
        # =====
        pass
```

3.2 Exploratory Data Analysis

```
# TODO: Analyze the simulated learner data

def analyze_learner_data(env, num_episodes=1000):
    """
    Collect episodes using random policy and analyze:
    - Distribution of episode lengths
    - Reward distribution
    - Completion rates by difficulty
    - Knowledge gain distribution
    """
    # ===== TODO =====
    # 1. Collect episodes with random actions
    # 2. Plot episode length histogram
    # 3. Plot reward distribution
    # 4. Plot completion rate vs module difficulty
    # 5. Compute summary statistics
    # =====
    pass
```

3.3 Baseline Model

```
# TODO: Implement the rule-based baseline recommender

class RuleBasedRecommender:
    """
    Baseline recommender that matches modules to learner level.
    Always picks the module closest to the learner's current mastery level.
    """
    def recommend(self, state, candidate_modules):
        """
        Args:
            state: learner state vector
            candidate_modules: list of module indices

        Returns:
            action: index of recommended module
        """
        # ===== TODO =====
        # 1. Extract learner mastery from state
        # 2. Find module with closest difficulty match
        # 3. Return module index (no exploration)
        # =====
        pass
```

3.4 Policy Gradient Model Architecture

```
# TODO: Implement the Actor-Critic content recommender

class ContentRecommenderActor(nn.Module):
    """
    Policy network for content recommendation.
    Maps learner state to a distribution over candidate modules.
    """
    def __init__(self, state_dim, hidden_dim=256, num_candidates=100):
        """
        Architecture:
        - State encoder: state_dim -> hidden_dim -> hidden_dim
        - Module scorer: takes encoded state + module embedding, outputs scalar score
        - Softmax over candidate scores
        """
        super().__init__()
        # ===== TODO =====
        # Build the actor network
        # =====
        pass

    def forward(self, state, candidate_embeddings):
        """
```

```

Args:
    state: (batch, state_dim)
    candidate_embeddings: (batch, K, embed_dim)

Returns:
    action_probs: (batch, K) probability distribution over candidates
"""

# ===== TODO =====
pass

class ContentRecommenderCritic(nn.Module):
    """Value network: estimates expected session return from current state."""
    def __init__(self, state_dim, hidden_dim=256):
        super().__init__()
        # ===== TODO =====
        pass

    def forward(self, state):
        # ===== TODO =====
        pass

```

3.5 Training Loop

```

# TODO: Implement the Actor-Critic training loop

def train_content_recommender(env, actor, critic, num_episodes=5000,
                               lr_actor=1e-3, lr_critic=1e-3,
                               gamma=0.95, entropy_coeff=0.01):
    """
    Train the content recommender using Actor-Critic with entropy bonus.

    Returns:
        reward_history, completion_rate_history, exploration_history
    """

    # ===== TODO =====
    # 1. Initialize optimizers
    # 2. For each episode:
    #     a. Reset environment
    #     b. Collect full episode (states, actions, rewards, log_probs)
    #     c. Compute returns G_t
    #     d. Compute advantages A_t = G_t - V(s_t)
    #     e. Compute actor loss, critic loss, entropy bonus
    #     f. Update both networks
    # 3. Track metrics: rewards, completion rates, exploration ratio
    # =====
    pass

```

3.6 Evaluation

```

# TODO: Evaluate the trained recommender against baseline

def evaluate_recommender(env, model, num_episodes=500):
    """
    Evaluate without exploration (greedy policy).

    Returns:
        completion_rate, avg_session_return, avg_knowledge_gain, unique_modules_recommended
    """

    # ===== TODO =====
    # 1. Run episodes with greedy action selection
    # 2. Compute all evaluation metrics
    # 3. Return metrics dictionary
    # =====
    pass

```

3.7 Error Analysis

```
# TODO: Analyze failure modes
```

```

def error_analysis(env, model, num_episodes=200):
    """
    Identify when and why the recommender fails.

    Analyze:
    - Which learner profiles get low rewards
    - Which modules are over/under-recommended
    - Session length distribution for completed vs abandoned sessions
    - Knowledge gain vs session return correlation
    """
    # ===== TODO =====
    pass

```

3.8 Deployment Simulation

```

# TODO: Simulate A/B test deployment

def simulate_ab_test(env, baseline_model, pg_model, num_users=1000, sessions_per_user=5):
    """
    Simulate an A/B test between rule-based and policy gradient recommenders.

    Args:
        env: learner environment
        baseline_model: rule-based recommender
        pg_model: trained policy gradient recommender
        num_users: number of simulated users
        sessions_per_user: sessions per user

    Returns:
        A/B test results with statistical significance
    """
    # ===== TODO =====
    # 1. Split users 50/50
    # 2. Run sessions for each group
    # 3. Compute metrics for each group
    # 4. Run statistical significance test (t-test)
    # 5. Report lift and p-value
    # =====
    pass

```

3.9 Ethics and Fairness

```

# TODO: Analyze fairness across learner demographics

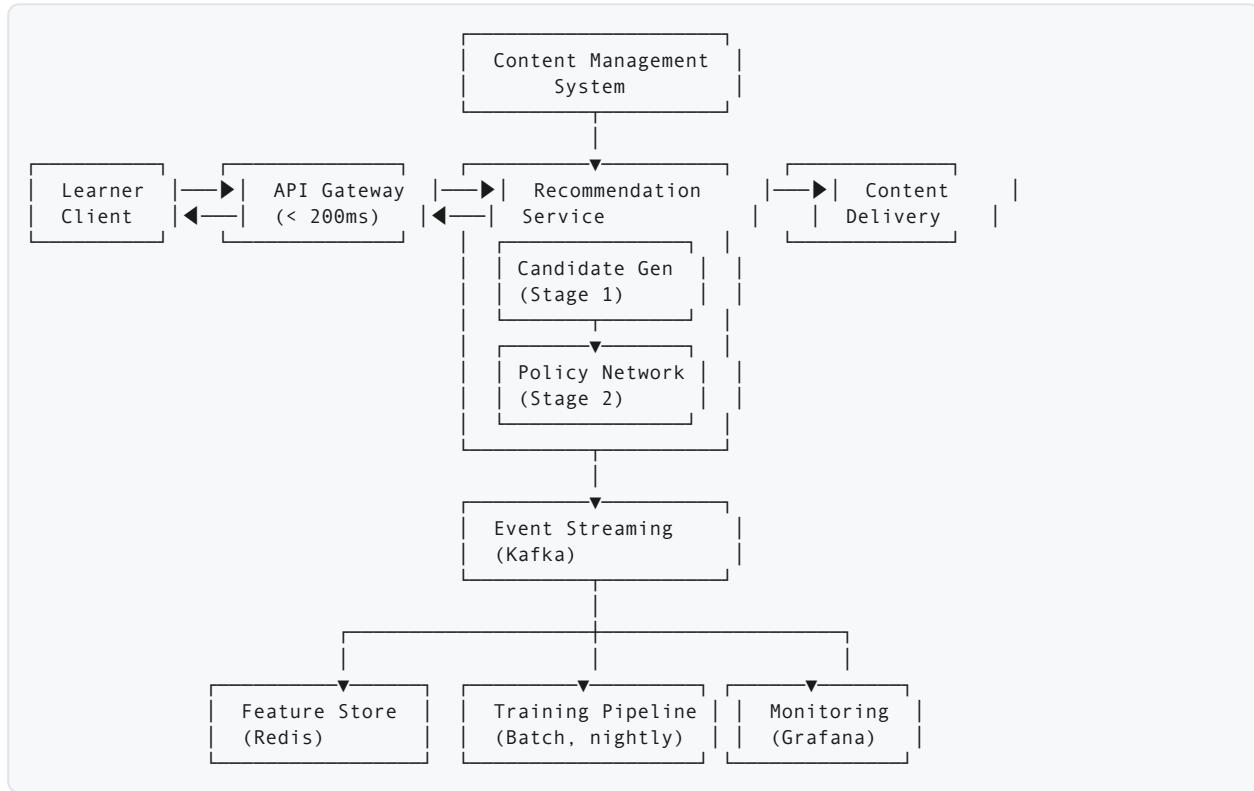
def fairness_analysis(env, model, num_episodes=1000):
    """
    Check if the recommender performs equally well across:
    - Different initial knowledge levels (beginner, intermediate, advanced)
    - Different learning speeds (fast, medium, slow)
    - Different session lengths (short, medium, long)

    Report any disparities.
    """
    # ===== TODO =====
    # 1. Create learner cohorts by knowledge level
    # 2. Run episodes for each cohort
    # 3. Compare completion rates and knowledge gains
    # 4. Flag any cohort with >10% performance gap from the best
    # =====
    pass

```

Section 4: Production and System Design Extension

System Architecture



API Design

```
POST /api/v1/recommend
Request:
{
  "user_id": "u_12345",
  "session_id": "s_67890",
  "context": {
    "device": "mobile",
    "time_of_day": "evening",
    "session_duration_so_far_seconds": 420
  }
}

Response (< 200ms):
{
  "module_id": "m_789",
  "module_title": "Introduction to Linear Regression",
  "confidence": 0.72,
  "alternatives": [
    {"module_id": "m_456", "confidence": 0.18},
    {"module_id": "m_123", "confidence": 0.10}
  ],
  "explanation": "Selected based on your recent progress in statistics fundamentals."
}
```

Model Serving

- **Inference framework:** ONNX Runtime for model inference (sub-10ms)
- **Candidate generation:** Pre-computed nearest-neighbor index (FAISS) for top-100 retrieval (sub-5ms)

- **Feature computation:** Redis feature store with pre-computed learner embeddings (sub-2ms)
- **Total latency budget:** < 50ms (well within the 200ms API requirement)
- **Throughput:** 10,000 recommendations/second on a single GPU instance

Monitoring and Drift Detection

Real-time metrics (Grafana dashboards): - Recommendation latency (P50, P95, P99) - Session completion rate (rolling 1-hour window) - Average session return (rolling 1-hour window) - Policy entropy (should not collapse below threshold)

Drift detection: - Feature drift: Monitor distribution of learner state vectors (KL divergence) - Reward drift: Detect shifts in completion rate or engagement signals - Action drift: Monitor whether the policy concentrates on too few modules

Alert thresholds: - Completion rate drops below 40% for > 1 hour - Policy entropy drops below 1.0 (policy collapse) - P99 latency exceeds 150ms - Any single module recommended > 5% of the time

A/B Testing Framework

- **Traffic split:** 90% current model / 10% new model initially
- **Ramp-up schedule:** 10% -> 25% -> 50% -> 100% over 4 weeks
- **Primary metric:** 7-day learner completion rate
- **Guardrail metrics:** Session duration, skip rate, NPS score
- **Statistical framework:** Sequential testing with alpha-spending function to enable early stopping

CI/CD Pipeline

1. **Data validation:** Check for feature drift in training data
2. **Model training:** Nightly retraining on the last 7 days of interaction data
3. **Offline evaluation:** Compare new model against current production model on held-out sessions
4. **Shadow deployment:** Run new model in shadow mode, compare recommendations without serving them
5. **Canary release:** Serve to 1% of traffic, monitor all metrics
6. **Full rollout:** Gradual traffic ramp-up with automatic rollback on metric degradation

Cost Estimation

Component	Monthly Cost
GPU inference (2x A10G)	\\$2,400
Redis feature store	\\$800

Component	Monthly Cost
Kafka event streaming	\\$600
Training compute (nightly)	\\$1,200
Monitoring infrastructure	\\$400
Total	\\$5,400/month

At 2.3M users and \\$18/user/month, infrastructure cost is 0.013% of revenue — negligible.

References

1. Sutton, R. S., McAllester, D., Singh, S., & Mansour, Y. (1999). "Policy Gradient Methods for Reinforcement Learning with Function Approximation." NeurIPS.
2. Williams, R. J. (1992). "Simple Statistical Gradient-Following Algorithms for Connectionist Reinforcement Learning." Machine Learning.
3. Ie, E., et al. (2019). "SlateQ: A Tractable Decomposition for Reinforcement Learning with Recommendation Sets." AAAI.
4. Chen, M., et al. (2019). "Top-K Off-Policy Correction for a REINFORCE Recommender System." WSDM.